

Hanjun Luo

Abu Dhabi, UAE — hl6266@nyu.edu — [Google Scholar](#) — [GitHub](#) — [Homepage](#)

EDUCATION

Ph.D. in Computer Science

New York University

Expected completion: 2030

Bachelor of Engineering in Computer Engineering

Zhejiang University

Year of completion: 2025

Bachelor of Science in Computer Engineering

University of Illinois Urbana-Champaign

Year of completion: 2025

RESEARCH INTERESTS

- Evaluation of LLMs, MLLMs, and LLM agents
- LLM/agent-as-a-judge methodology
- Human-agent collaboration
- Safety of LLM agents
- Bias and trustworthiness in generative models
- LLM-based multilingual information extraction

PUBLICATIONS

Luo, H., Huang, Z., Chung, S., Wang, Y., Jin, Y., Li, J., Li, J., Li, X., & Salam, H. (2026). AtelierEval: Agentic Evaluation of Humans & LLMs as Text-to-Image Prompts. *International Conference on Machine Learning (ICML 2026)*.

Luo, H., Ni, C., Wen, J., Huang, Z., Wang, Y., Liao, B., Chung, S., Jin, Y., Li, J., Li, X., Xu, W., Wang, X., & Salam, H. (2026). CentaurEval: Benchmarking Human-in-the-Loop Value in Agentic Coding. *International Conference on Machine Learning (ICML 2026)*.

Luo, H., Liu, Q., Zhang, J., Pang, H., Wen, J., Ma, Y., Yao, Y., Zhang, C., Zhang, H., Chen, Y., & Salam, H. (2026). AutoCRAT: Within-trajectory joint control of stochasticity and compute for LLM reasoning. *Conference on Empirical Methods in Natural Language Processing (EMNLP 2026 Findings)*.

Luo, H., Huang, Z., Huang, H., Deng, Z., Chen, R., Li, X., Liu, Z., & Salam, H. (2026). BiasIG: Benchmarking multi-dimensional social biases in text-to-image models. *International Joint Conference on Neural Networks (IJCNN 2026, Oral)*.

Luo, H., Dai, S., Ni, C., Li, X., Zhang, G., Wang, K., Liu, T., & Salam, H. (2025). AgentAuditor: Human-level safety and security evaluation for LLM agents. *Conference on Neural Information Processing Systems (NeurIPS 2025)*, poster.

Luo, H., Jin, Y., Li, X., Liu, X., Chen, R., Shang, T., Wang, K., Wen, Q., & Liu, Z. (2025). DynamicNER: A dynamic, multilingual, and fine-grained dataset for LLM-based named entity recognition. *Conference on Empirical Methods in Natural Language Processing (EMNLP 2025)*.

Luo, H., Deng, Z., Chen, R., & Liu, Z. (2024). FAIntbench: A holistic and precise benchmark for bias evaluation in text-to-image models. *ICML 2024 Workshop on Data-centric Machine Learning Research (DMLR)*.

Jin, Y., Wang, J., Chai, H., **Luo, H.**, Hu, L., Li, X., Zheng, H., Hu, P., Ye, Q., & Hu, H. (2027). MobileProbe: Adaptive Security Evaluation of Mobile Agents against Trajectory-Aware Injection. *Network and Distributed System Security Symposium (NDSS 2027)*.

Li, J., Chen, Z., **Luo, H.**, & Salam, H. (2026). Prefix: Understand and adapt to user preference in human-agent interaction. *Annual Meeting of the Association for Computational Linguistics (ACL 2026 Findings)*.

Jin, Y., Du, X., **Luo, H.**, Wang, Z., Hu, H., Wang, X., & Li, X. (2026). AudioStealer: Extracting audio prompts via Shapley value-guided query search. *Annual Meeting of the Association for Computational Linguistics (ACL 2026 Findings)*.

Li, K., Shen, C., Liu, Y., Han, J., Zheng, K., Zou, X., Wang, Z., Du, X., Zhang, S., **Luo, H.**, et al. (2026). AudioTrust: Benchmarking the multifaceted trustworthiness of audio large language models. *International Conference on Learning Representations (ICLR 2026)*, poster.

Sun, M., Zhao, Z., Chai, W., **Luo, H.**, Cao, S., Zhang, Y., Hwang, J.-N., & Wang, G. (2024). UniAP: Towards universal animal perception in vision via few-shot learning. *AAAI Conference on Artificial Intelligence (AAAI 2024)*.

RESEARCH EXPERIENCE

LLM/MLLM/Agent Evaluation

- AtelierEval (ICML 2026): Build a benchmark with 360 expert-crafted tasks to evaluate text-to-image prompting proficiency for both MLLMs and human prompters.
- CentaurEval (ICML 2026): Build a coding benchmark for both LLM agents and human developers with 45 collaboration-necessary templates and 450 instantiated tasks, showing low pass rates for LLM-only (0.67%) and human-only (18.89%) settings but a clear collaborative gain to 31.11%.
- AgentAuditor (NeurIPS 2025): Build ASSEBench, an LLM-as-a-Judge benchmark with 2,293 annotated interaction records covering 15 risk types across 29 application scenarios for end-to-end agent safety/security evaluation.
- DynamicNER (EMNLP 2025): Build a dynamic and multilingual benchmark spanning 8 languages, 155 entity types, and 3 granularity levels to evaluate robustness and generalization of LLM-based NER.
- Prefix (ACL Findings 2026): Introduce a configurable benchmark with Interaction-as-a-Tool (IaaT), 31 preference settings, and 14 attributes to evaluate LLM-agent preference adaptation beyond task accuracy.
- AudioTrust (ICLR 2026): Construct a multifaceted benchmark spanning 6 trustworthiness dimensions, 26 sub-tasks, and 4,420+ audio samples, with evaluation across 14 SOTA ALLMs under 18 experimental configurations.

LLM/Agent-as-a-Judge

- AgentAuditor (NeurIPS 2025): Introduce a novel Agent-as-a-Judge method with memory-augmented reasoning for agent safety and security evaluation.
- BiasIG (IJCNN 2026 Oral): Construct a fine-tuned MLLM-based judge for bias in text-to-image models.
- Prefix (ACL Findings 2026): Propose a composite LLM-as-a-Judge protocol over seven UX dimensions for scalable preference-alignment scoring in agent interactions.

LLM Agent Safety

- AgentAuditor (NeurIPS 2025): Introduce both an Agent-as-a-Judge method and a benchmark for agent safety and security evaluation.
- A Comprehensive Survey in LLM(-agent) Full Stack Safety: Lead the benchmarking component and systematize LLM and agent safety risks and mitigation across data, training, and deployment.

Human-Agent Collaboration

- CentaurEval (ICML 2026): Design collaboration-necessary tasks to quantify human-in-the-loop value in coding and verify them through human studies and LLM benchmarking.
- Prefix (ACL Findings 2026): Propose a preference-aware LLM-agent collaboration method that infers implicit user preferences and adapts confirmation style and pacing in multi-turn workflows.

Bias Evaluation and Debiasing in Text-to-Image Models

- BiasIG (IJCNN 2026 Oral): Construct a multi-dimensional social-bias benchmark with 47,040 prompts across 4 dimensions and benchmark 8 T2I models with 3 debiasing methods.
- FAIntbench (ICML-DMLR 2024): Build a holistic benchmark for fine-grained text-to-image bias evaluation.
- VersusDebias: Propose a zero-shot debiasing framework and analyze trade-offs.

Audio Model Trustworthiness

- AudioTrust (ICLR 2026): Based on the benchmark, identify key findings on audio-specific vulnerabilities from non-semantic cues (e.g., accent, timbre, background noise) across 14 ALLMs and 18 evaluation settings.
- AudioStealer (ACL Findings 2026): Design a two-stage Shapley-guided black-box prompt extraction attack. Introduce Prompt2Music, a prompt-stealing benchmark with 5,000 prompt-audio pairs generated from 4 text-to-music models.

LLM-based Multilingual Information Extraction

- DynamicNER (EMNLP 2025): Develop a realistic multilingual setting for NER, covering evolving entities and fine-grained categories.

VISITING EXPERIENCE

Tsinghua University

2026.5–2026.8

Visiting Researcher, Prof. Yipeng Dong

TEACHING EXPERIENCE

CS101: Intro to Computing

2024.9–2024.12

Teaching Assistant, Zhejiang University, Prof. Wee Liat Ong

ADDITIONAL SKILLS

- **Languages:** Chinese (native), English (fluent).
- **Programming:** Python, SQL, LaTeX.
- **Photography:** Signed contributor for Visual China Group (VCG); experienced in landscape photography.